

Cloud Provider Sustainability Current Status and Future Directions

DevSusOps: Adding sustainability concerns to development and operations

Adrian Cockcroft - OrionX.net

@adrianco@mastodon.social

@DevSusOps@hachyderm.io

October 2023

Agenda

Why does sustainability matter?
Why is this hard?

Terminology and mental models
Cloud provider similarities and differences
Cloud carbon measurements
Things developers need to build

The background image shows a wide-angle view of a coastal landscape at sunset. In the foreground, there's a body of water with small waves. Across the water, a town with numerous houses and buildings is nestled at the base of several large, rugged mountains. The sky is filled with dramatic, billowing clouds that are illuminated from behind by the setting sun, creating a warm, golden glow. The overall atmosphere is serene and beautiful.

Apology

There is a lot of dense content in this talk

also I don't have a graphic design budget
to make it look nicer nowadays...

Leave the world habitable for future generations

Market transition risks

Regulatory compliance

Physical risks to business assets

Why does sustainability matter?

“Green” market positioning

Employee enthusiasm

Reduced costs now or in the future

Social license to operate

Regulators and market pressure...

Companies need to report carbon emissions, assess climate risk and pay border taxes on carbon

California SB-253 Climate Corporate Data Accountability Act
https://leginfo.legislature.ca.gov/faces/billTextClient.xhtml?bill_id=202320240SB253

California SB-261 Greenhouse gases: Climate-related financial risk
https://leginfo.legislature.ca.gov/faces/billNavClient.xhtml?bill_id=202320240SB261

EU Carbon Border Tax
<https://www.politico.eu/article/europe-climate-fight-global-carbon-border-adjustment-mechanism-cbam-tax/>

EU Corporate Sustainability Reporting Directive (CSRD)
https://finance.ec.europa.eu/capital-markets-union-and-financial-markets/company-reporting-and-auditing/company-reporting/corporate-sustainability-reporting_en

SEC Proposes Rules to Enhance and Standardize Climate-Related Disclosures for Investors
<https://www.sec.gov/news/press-release/2022-46>

Things we can do now...

Development

- Optimize code
- Choose faster languages and runtimes
- Efficient algorithms
- Faster implementations
- Reduce logging
- Reduce retry storms and work amplification

Operations

- Higher utilization
- Automation
- Relax over-specified requirements
- Archive and delete data sooner
- Deduplicate data
- Choose times and locations carefully

A wide-angle photograph of a tropical coastal area. In the foreground, there's a body of water with small waves. Along the shore, there's a cluster of houses and buildings, some with white roofs and others with darker ones. Several tall palm trees are scattered throughout the scene. In the middle ground, there's a flat, agricultural landscape with various fields and some utility poles. The background is dominated by a range of mountains covered in lush green vegetation. One prominent mountain on the right has a distinct reddish-brown patch of soil or rock on its side.

Measuring Carbon

Why is this hard?

We just have to multiply two numbers, right?


Energy * Carbon Content

Scopes of Carbon

(Using some slides I made for AWS)

Scope 1—fuel consumed

COUNTED BY WHOEVER OWNS THE FUEL WHEN IT BURNS



Car




Gas pump

Fireplace

Furnace

Gas meter




Cooking

Electrify everything to take Scope 1 to zero

Scope 2—energy used


ELECTRICITY USE IS COUNTED ONLY ONCE WHERE IT IS CONSUMED



Windmill
Power station
Grid mix



Heat pump Solar panels Batteries Electric car




Induction range
(gas ranges are also a big cause of indoor pollution that leads to asthma)

Change grid mix to renewable power and store renewable energy in batteries to reduce Scope 2

Scope 3—everything else


SUPPLY CHAIN AND INVESTMENTS



Scope 3 depends a lot on what kind of business you are in

The carbon footprint of datacenters


TOP-DOWN CONTEXT




© 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

The carbon footprint of datacenters

TOP-DOWN CONTEXT



POWER PYLONS



DELIVERING NEW RACKS AND RECYCLING OLD RACKS


© 2021, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Your suppliers need to report scope 1,2,3 to you

You need to report scope 1,2,3 to your customers

Greenhouse Gas Protocol - <https://ghgprotocol.org>

Scopes depend who you are...



<https://learn.microsoft.com/en-us/industry/sustainability/api-calculation-method> - minor update 6/30/2023

As the carbon content of energy tends to zero
scope 3 dominates carbon footprints

Energy * Carbon Content + Supply Chain

In the EU and US cloud regions this happened a few years ago

This data does NOT include all the private power generation that cloud providers use to clean up the local grids and meet their goals


Data Centers Impact distribution by country

Region	gCO2e/KWh	Scope 1	Scope 2	Scope 3	Scope 1	Scope 2	Scope 3	TOTAL	var Vs EU-28
EU-28	265	0.7%	68.2%	31.1%	0.7	68.2	31.1	100.0	
Sweden	27	1.8%	19.7%	80.3%	0.7	6.9	31.1	38.7	-61%
Finland	148	1.0%	55.5%	44.5%	0.7	38.1	31.1	69.9	-30%
Norway	37	1.7%	24.7%	75.3%	0.7	9.5	31.1	41.3	-59%
France	85	1.3%	42.1%	57.9%	0.7	21.9	31.1	53.7	-46%
Poland	780	0.3%	86.6%	13.4%	0.7	200.7	31.1	232.5	133%
Ireland	342	0.6%	74.0%	26.0%	0.7	88.0	31.1	119.8	20%

gCO2e/KWh: source electricitymaps.org and European Environment Agency for EU-28 data

Scope 1, 2 & 3 of EU-28 from "Digital technologies in Europe: an environmental life cycle approach" Dec. 2021

Get more recent data from electricitymaps.com (not .org) Thanks to Cyril Deblois - AWS Sustainability Lead - Cloud Economics for this chart.



The world's carbon intensity viewed on electricitymaps.com

Cloud Provider Scope 3 Differences

Amazon Web Services (AWS) - Still no public scope 3 (they say they are working on it)

- Discussion of why not on Adrian's Medium Blog LinkedIn page and a Computer Weekly story
- <https://www.computerweekly.com/news/365531874/Amazon-denies-claims-hiring-freeze-is-slowing-AWS-sustainability-work>
- Summary: work continues, slowly, don't hold your breath... escalate to get rough estimates under NDA if needed


Microsoft Azure - Detailed scope 3 data and API

- Launched in 2021 - excellent methodology white paper - <https://go.microsoft.com/fwlink/?linkid=2161861>
- Includes recycling, doesn't include transport and buildings yet
- Numbers will increase a bit when these are added...


Google Cloud Platform (GCP) - Detailed scope 3

- Includes transport and buildings
- Doesn't include end of life recycling (estimated to be immaterial)

Scope 3 Methodology



<https://learn.microsoft.com/en-us/industry/sustainability/api-calculation-method>



How to reduce Scope 3

- Use less hardware
- Increase utilization
- Use hardware you own for longer before replacing it
- Use less of the latest hardware - faster and more efficient
- ARM architecture like AWS Graviton uses less power than Intel/AMD
- Use fewer bigger systems - per-system carbon overhead is large
- Select larger underlying system types for instance VMs/Nodes

Transportation contribution to carbon is very small, so scope 3 is basically the same everywhere in the world

List of carbon info for hardware: <https://github.com/Boavizta/environmental-footprint-data/blob/main/boavizta-data-us.csv>
 For example: https://i.dell.com/sites/csdocuments/CorpComm_Docs/en/carbon-footprint-poweredge-r630.pdf

OK, but scope 3 is about hardware, so it's an “Ops problem”, we developers just have to multiply two numbers, right?

Energy * Carbon Content

Energy

How much energy does a specific line of code use?

How much energy does a transaction use? (SCI)

How much energy does a specific workload use?

How much energy does a container use?

How much energy does a specific cloud instance use?

Cloud providers don't tell you any of these things...

See <https://www.noureddine.org/research/joular/joularjx> for real-time JVM energy monitoring

See <https://github.com/sustainable-computing-io/kepler> for Kubernetes energy analysis

Measuring Energy Usage

Decided to figure out how to measure the energy used by a desktop computer and see if I can figure out a way to identify different workloads.

I have a Mac Studio M1 to run the workload, and an old MacBook laptop to run data collection on, so that it doesn't add to the workload.

First thing we need is a power monitoring plug that has an API. The TP-Link Kasa platform seems like a good place to start. It has a python based API available [on GitHub](#).

```
% pip3 install python-kasa
```

I ordered a [Kasa KP115 smart plug from Amazon](#) for \$22.99.



DIY Energy Usage Measurement

Repeatability and Reproducibility challenges

Lookup how to do Gage R&R statistical analysis!

Variance from one test run to the next
Variance from one system to another
Lots of confounding effects
Activity monitor uses the most power...
Possibly a fun project to try at home?

Datacenter systems can use Redfish
<https://www.dmtf.org/standards/redfish>

Activity Monitor		Applications in last 12 hours	
App Name	Energy Imp...	12 hr Power	
Activity Monitor	34.0	5.19	
Milkshake	16.1	0.43	
Safari	0.1	44.22	
X-AIR-Edit	10.8	2.18	
Discord	0.1	5.73	
Amazon Chime	2.5	0.77	
Dropbox	2.1	9.85	
Photos	1.0	0.44	
zoom.us	0.9	36.42	
Slack	0.0	4.11	
In Your Face	0.5	0.12	
LogTune	0.4	-	
Finder	0.1	2.12	
Mona	0.0	0.67	

<https://dev.to/adrianco/measuring-energy-usage-5ip>

OK, Energy is a pain to measure...
Where do we get the Carbon Content from?

Energy * Carbon Content

Cloud Carbon Content - It depends...

What is the grid mix for energy generation in that location?

When was the grid mix measured? (Real time data isn't free)

Do you have hourly grid mix? (GCP does but doesn't share it)

Do the grid mix numbers for a specific hour change over time? (yes)

How much private cloud provider energy was used? (they don't say)

How much "Bundled REC" energy was used? (they don't say)

How much "Unbundled REC" energy was used? (they don't say)

A wide-angle photograph of a bright rainbow stretching from the upper left towards the lower right against a clear, pale blue sky. Below the sky, a lush, green tropical landscape is visible, featuring several palm trees and a distant, low-lying shoreline under a hazy horizon.

How can we reduce the carbon emissions
of the systems we build?

Saving carbon is directionally the same
as saving \$\$\$

(In almost all cases)

But Carbon/\$ varies a lot...

A wide-angle landscape photograph showing a range of mountains in the background. The mountains are covered in lush green vegetation, with some areas appearing darker and more rugged. In the middle ground, there's a flat, lighter-colored area that looks like a coastal plain or a valley floor. The foreground is a dark, calm body of water, likely an ocean or a large lake. The sky above is a clear, vibrant blue, dotted with wispy, white, cumulus-like clouds.

Focus on how to measure carbon

We're used to reporting and optimizing
throughput, latency, utilization, capacity, cost...


Carbon is just another metric for monitoring tools to report

Measuring carbon emitted by a workload
How hard can it be?

Power Usage Efficiency (PUE) accounts for losses and cooling overhead
but value depends on where it is measured

Power distribution and UPS batteries may be centralized or distributed in
different datacenter architectures


Grid Mix is obtained
from the local utility bill,
usually a month or so in
arrears, combined with
any power purchase
agreements (PPAs) or
renewable energy
credit (REC) purchases



Scope 2 Carbon = Power Mix * PUE * Capacity Used * Emissions factor per capacity

Two methodologies for reporting carbon
Location based vs. Market based

Market vs. Location Reporting for Scope 2



Why does it matter?

AWS and Azure carbon data is Regional Market based
e.g. PPA generation is connected to the same grid 😊

Google's 100% renewable since 2017 claim is Global Market based
e.g. Generation in USA and EU was counted against Singapore 😞

Google's current data, API and 24/7 work is Location Based 🎉
Numbers can't be compared, but more useful for tuning work

As the utility grid decarbonizes over time, it matters less...

Problems...

Scope 2 Carbon = Power Mix * PUE * Capacity Used * Emissions factor per capacity



Utility bills delayed by a month or more, depending where in the world they are...

PPA - Power Purchase Agreements are contracts to build and consume power generation capacity e.g. Amazon has over 20GW of PPAs (Solar, Wind, Battery)

REC - Renewable Energy Credits are purchases of renewable generation capacity on the open market, from existing generators, the energy can only be claimed once, but can be claimed later

RECs may be used to “top up” on top of PPAs, helps fund existing renewable capacity, but don’t create additional capacity, like PPAs do

Good RECs and Bad RECs

Guarantee of Origin (in EU) and Renewable Energy Credit (in USA)

- Similar but have some detailed differences in how they are calculated, regulated and traded

Local Market REC

- Reserve the renewable energy from a supplier on the same grid - GOOD!

Cross-border/non-local market REC

- Renewable energy that can't flow to you
- Extra funding helps to subsidize existing renewable generation - GOOD
- Argument that carbon reduced in one part of the world offsets other parts of the world...
- Often used as a cheap carbon offset - greenwashing - BAD

<https://www.ecocostsvalue.com/lca/gos-and-recs-in-lca/>

<https://sustainability.google/progress/projects/ppa/>

<https://www.blog.google/outreach-initiatives/environment/meeting-our-match-buying-100-percent-renewable-energy/>

Problems...

Scope 2 Carbon = Power Mix * PUE * Capacity Used * Emissions factor per capacity



Grid mix changes every month but is mostly getting better over time

Hourly data for 24/7 grid mix starting to appear for some suppliers in some parts of the world


GCP and Azure publicly working towards 24/7

cloudcarbonfootprint.org open source tool estimates don't include PPAs or RECs

Also see [FlexiDAO.com](#) startup working on 24/7 energy trading

Power Mix Problems and Misconceptions

Scope 2 Carbon = Power Mix * PUE * Capacity Used * Emissions factor per capacity



Misconceptions - most people don't understand that RECs and PPAs are excluded from the grid mix

PPAs are private and don't affect the grid mix for everyone else


Use of local market RECs is good, but makes the grid mix worse for everyone

RECs can be traded for up to a year, so you can pay to offset past emissions, and that changes the grid mix for the month, for up to a year later...

<https://www.ecocostsvalue.com/lca/gos-and-recs-in-lca/>

Power Mix Problems and Misconceptions

Scope 2 Carbon = Power Mix * PUE * Capacity Used * Emissions factor per capacity



Misconceptions - most people don't understand that RECs and PPAs are excluded from the grid mix

PPAs are private and don't affect the grid mix for everyone else

Use of local market RECs is good, but makes the grid mix worse for everyone

RECs can be traded for up to a year, so you can pay to offset past emissions, and that changes the grid mix for the month, for up to a year later...

<https://www.ecocostsvalue.com/lca/gos-and-recs-in-lca/>

Problems...

Scope 2 Carbon = Power Mix * PUE * Capacity Used * Emissions factor per capacity



Power Usage Efficiency is not well standardized. Values can be compared for similar datacenter designs, but some good efficiency improvements may move components from the infrastructure to the equipment, and make PUE worse.

AWS doesn't publish PUE and points to this blog

<https://perspectives.mvdirona.com/2009/06/pue-and-total-power-usage-efficiency-tpue/>

Azure published PUE

<https://azure.microsoft.com/en-us/blog/how-microsoft-measures-datacenter-water-and-energy-use-to-improve-azure-cloud-sustainability/>

Google published PUE

<https://www.google.com/about/datacenters/efficiency/>

Problems...

Scope 2 Carbon = Power Mix * PUE * Capacity Used * Emissions factor per capacity



Dedicated compute and storage capacity is relatively easy to account for, but sliced instances, shared services, network equipment etc. Is a challenge to allocate

If you are operating a multi-tenant service and want to report carbon, then you have a problem figuring out how much capacity to allocate to each customer, and how much is overhead that you should own yourself...

Problems...

Scope 2 Carbon = Power Mix * PUE * Capacity Used * Emissions factor per capacity




Need to know how much power each instance type, storage class or service uses, where that capacity is located, and how much carbon is emitted per kWh

This depends on utilization and other overheads confuse measurements

Mostly unavailable data at present, although boavizta.org has some estimates

Takeaways



What can you do today?

Use Less Stuff

This will make the biggest difference

AWS "Better than 95% Renewable 2021" 13 Regions


To achieve our goal of powering our operations with 100% renewable energy by 2025—five years ahead of our original 2030 target—Amazon contracts for renewable power from utility scale wind and solar projects that add clean energy to the grid. These new renewable projects support hundreds of jobs while providing hundreds of millions of dollars of investment in local communities. We also may choose to support these grids through the purchase of environmental attributes, like Renewable Energy Certificates and Guarantees of Origin, in line with our [Renewable Energy Methodology](#).

As a result, in 2021, the following AWS Regions were powered by over 95% renewable energy:

- US East (Northern Virginia)
- GovCloud (US-East)
- US East (Ohio)
- US West (Oregon)
- GovCloud (US-West)
- US West (Northern California)
- Canada (Central)
- Europe (Ireland)
- Europe (Frankfurt)
- Europe (London)
- Europe (Milan)
- Europe (Paris)
- Europe (Stockholm)

Mostly located in the USA and Europe but more power projects in Asia than Azure and GCP

100% Renewable by 2025 on an annual market basis commitment for Amazon (not just AWS)



<https://sustainability.aboutamazon.com/environment/the-cloud?energyType=true>


AWS "100% Renewable 2022" 19 Regions

We contract for renewable power from utility scale wind and solar projects that add clean energy to the grid. These new renewable energy projects support hundreds of jobs while providing hundreds of millions of dollars of investment in local communities. We also may choose to support these grids through the purchase of environmental attributes, like Renewable Energy Certificates and Guarantees of Origin, in line with our [Renewable Energy Methodology](#).

In 2022, the electricity consumed in the following 19 AWS Regions was attributable to 100% renewable energy:

- U.S. East (Northern Virginia)
- GovCloud (U.S. East)
- U.S. East (Ohio)
- U.S. West (Oregon)
- GovCloud (U.S. West)
- U.S. West (Northern California)
- Canada (Central)
- Europe (Ireland)
- Europe (Frankfurt)
- Europe (London)
- Europe (Milan)
- Europe (Paris)
- Europe (Stockholm)
- Europe (Spain)
- Europe (Zurich)
- Asia-Pacific (Mumbai)
- Asia-Pacific (Hyderabad)
- China (Beijing)
- China (Ningxia)

Since 2021, US/Europe goes to 100%, India and China added... still leaves Japan, Korea, Indonesia, Australia, Brazil, South Africa to do.




<https://sustainability.aboutamazon.com/environment/the-cloud?energyType=true>

Azure Global Renewable Projects


Mostly located in the USA
and Europe today

100% Renewable on an
annual market basis by
2025 commitment




<https://infrastructuremap.microsoft.com/explore>

2021 GCP 24x7 Hourly Differences Worldwide



Data is based on hourly 24x7 location model, not directly comparable to annual market model.

2022 GCP 24x7 Hourly Differences Worldwide



Not much progress year on year, worse over-all slipping from 66% to 64%

24/7 Carbon Free by 2030 commitment

Use any cloud provider but try to minimize use of Asia regions for the next few years

There really isn't that much difference in carbon footprint between the cloud providers

They are all buying many gigawatts of renewables


All have the same grid challenges in Asia, although AWS is making more progress

Scope 3 is dominated by the same Asian chip suppliers

They are all much better than a typical enterprise datacenter

Measuring Carbon...

Compare APIs and Schemas across AWS, Azure and GCP - and propose what we really want



AWS Customer Carbon Footprint Tool (csv download)

Time: Monthly summary

Place: Continent level

Specificity: Account, EC2, S3, Other

Resolution: 0.001 Metric Tons of CO₂e

Models: Scope 1, market based scope 2

<https://aws.amazon.com/aws-cost-management/aws-customer-carbon-footprint-tool/>

<https://aws.amazon.com/blogs/aws/new-customer-carbon-footprint-tool/>

<https://aws.amazon.com/blogs/aws-cloud-financial-management/increased-visibility-of-your-carbon-emissions-data-with-aws-customer-carbon-footprint-tool/>

Sample queries for the EnrollmentEmission entity

Query type	Example
Emissions by enrollment	{serviceRoot}/emissions
Select certain fields	{serviceRoot}/emissions?\$select=enrollmentId,totalEmissions,scopeId
Include count	{serviceRoot}/emissions?\$count=true
Limit result count	{serviceRoot}/emissions?\$top=100
Paging	{serviceRoot}/emissions?\$skip=100&\$top=50
Filter by scope	{serviceRoot}/emissions?\$filter=ScopeId eq 1
Filter and aggregate	{serviceRoot}/emissions?\$apply=filter(Scopeld eq 1)/aggregate(\$count as Count, totalEmissions with average as Average, totalEmissions with sum as Sum)
Filter and group	{serviceRoot}/emissions?\$apply=filter(totalEmissions gt 0.05)/groupby((Scopeld), aggregate(\$count as Count))'

EnrollmentUsage entity

Represents a calculated usage factor of Microsoft cloud resources.

Property	Type	Notes
dateKey	int32	Date in yyyyymmdd format; dd is always 01.
enrollmentId	string	Also known as billing account ID.
orgName	string	
subscriptionId	string	
subscriptionName	string	
subService	string	For example, Azure Storage or Azure Compute.
azureRegionName	string	
usage	double	For more information about Microsoft's calculation methodology, go to Microsoft Cloud for Sustainability API calculation methodology .

Azure Carbon Footprint API Schema (OData Preview)

Time: Monthly summary

Place: Country and region specific

Specificity: Account, service

Resolution: 0.001 Metric Tons of CO2e

Models: Scope 1, market based scope 2, scope 3

<https://learn.microsoft.com/en-us/industry/sustainability/api-overview>

Carbon Footprint export data schema

After configuring an export, a BigQuery table with the following schema is created:

Field	Type	Description	Example
usage_month	Date	Month during which this usage occurred	2021-10-01
billing_account_id	String	The Cloud Billing account ID used for the carbon footprint export	007GD7-15C0E6-BDE6B3
project.id	String	Identifier of the Google Cloud project the carbon footprint is from	my-project
project.number	String	Number of the Google Cloud project the carbon footprint is from	882583874831
service.description	String	Readable name for the Google Cloud service the carbon footprint is from	Cloud Run
service.id	String	Identifier of the Google Cloud service the carbon footprint is from	F17B-412E-CB64
location.location	String	Location of the carbon footprint at the level of a country, region, or zone	us-central1-a
location.region	String	Cloud region for the carbon footprint. NULL if usage is multi-region or global	us-central1
carbon_model_version	Integer	Version of carbon model that produced this output. This value is updated whenever the model is changed	1
carbon_footprint_total_kgCO2e.after_offsets	Float	Total carbon footprint for the account, project, service, location, and month in kg of CO ₂ equivalent. Equivalent to scope 1 + scope 2 market-based + scope 3 + carbon offsets. Currently set to NULL	NULL
carbon_footprint_total_kgCO2e.market_based	Float	Total carbon footprint for the account, project, service, location, and month in kg of CO ₂ equivalent. Equivalent to scope 1 + scope 2 market-based + scope 3. Currently set to NULL	NULL
carbon_footprint_total_kgCO2e.location_based	Float	Total carbon footprint for the account, project, service, location, and month in kg of CO ₂ equivalent. Equivalent to scope 1 + scope 2 location-based + scope 3	1587.3
carbon_footprint_kgCO2e.scope1	Float	Total scope 1 carbon footprint for the account, project, service, location, and month in kg of CO ₂ equivalent	10
carbon_footprint_kgCO2e.scope2.location_based	Float	Total scope 2 location-based carbon footprint for the account, project, service, location, and month in kg of CO ₂ equivalent	400
carbon_footprint_kgCO2e.scope2.market_based	Float	Total scope 2 market-based carbon footprint for the account, project, service, location, and month in kg of CO ₂ equivalent. Currently set to NULL	NULL
carbon_footprint_kgCO2e.scope3	Float	Total scope 3 carbon footprint for the account, project, service, location, and month in kg of CO ₂ equivalent	986.8
carbon_offsets_kgCO2e	Float	Total carbon offsets for the account, project, service, location, and month in kg of CO ₂ equivalent. Currently set to NULL	NULL

Google Carbon Footprint BigQuery Export Schema

Time: Monthly summary

Place: Country, region and zone specific

Specificity: Account, project, service

Resolution: 0.1 Kilograms of CO₂e

Models: Scope 1, location based scope 2 with placeholders for market based, scope 3

<https://cloud.google.com/carbon-footprint/docs/data-schema>

Measuring Carbon

What do we want?

Real time carbon metrics for optimization
Just another metric like CPU Utilization

Reported via the same tools we already use
(e.g. from CloudWatch or Azure Monitor via Prometheus)

Real-Time Carbon Footprint Standard

A proposal - what @adrianco thinks would be useful...

Support: Same data for all cloud providers and datacenter automation tools

Time: Same resolution as existing monitoring tools, typically minutes

Place: Country, region and zone specific

Specificity: Account, project, service, instance, container, filesystem, etc...

Resolution: Grams of CO2e, energy in 0.001 joules (milliwatt-seconds)

Models: Scope 1, location *and* market based scope 2, scope 3

Energy data would be final, CO2e would be reported as mean and 95-percentile upper and lower confidence intervals, based on hourly updates to 24/7 renewable mix estimates including RECs and PPAs. CO2e could be re-processed to narrow the confidence interval to audit report quality after 1-2 months, and to finalize after REC settlement grid mix adjustments and supply chain scope 3 updates at 12 months.

<https://github.com/Green-Software-Foundation/real-time-cloud/blob/main/PRFAQ%20for%20RealTimeCarbonMetrics.md>


Real-Time Carbon Footprint Standard

Instance Energy - updates every minute (Cloud Provider)

Carbon Intensity - updates every hour (Cloud Provider)

CPU/RAM/Process Utilization/Energy - updates every minute (Kepler)


Kepler Allocation algorithm already estimates and allocates energy



GSF Real-Time Carbon Footprint Project

“All models are wrong, some models are useful, we want to make the cloud carbon footprint models less wrong and more useful, not perfect”

Project status, issues and meeting minutes are at: <https://github.com/Green-Software-Foundation/real-time-cloud>
 Miro board collecting information is at: https://miro.com/app/board/uXjVM1o59N4=/?share_link_id=388311040102



Comparison	Real-Time Carbon Footprint Standard	AWS Customer Carbon Footprint Tool	Azure Emissons Impact Dashboard	GCP Carbon Footprint
API Format	Kepler/Prometheus/ OpenTSDB	CSV Download	OData	BigQuery
Support	No-one yet...	AWS	Azure	GCP
Time	Minutes	Month	Month	Month
Place	Country, Region, Zone	Continent	Country, Region	Country, Region, Zone
Specificity	Accounts, Instances	Accounts, EC2, S3, Other	Account, Service	Account, Project, Service
Resolution	Grams CO2e +/- 0.001 Joules energy	0.001 Metric Tons CO2e	0.001 Metric Tons CO2e	0.1 Kg CO2e
Models	Scopes 1, 2L, 2M, 3	Scopes 1, 2M	Scopes 1, 2M, 3	Scopes 1, 2L, 3
Data Availability	Immediate, 2 month, 12 month final	3 month	2 month	2 month

Measuring Carbon

What tools could we build?


Re-use cost and performance optimization tooling for CO2e optimization

Carbon peak flattening scheduler, move work to other times and places

SaaS provider attribution and allocation tools for per-customer carbon reports

Architecture planning tools - e.g. what's the carbon footprint difference between running your own Kafka on EC2 vs. AWS MSK vs. Confluent vs. RedPanda vs. AWS Kinesis?

Your ideas?




Pathfinder API for supply chain carbon data
by the World Business Council for
Sustainable Development is an emerging
interchange protocol standard

<https://www.wbcsd.org/Programs/Climate-and-Energy/Climate/SOS-1.5/Resources/Pathfinder-Framework-Version-2.0>

One more thing...


Measuring Climate Impact Risk - to you!

[OS-Climate.org](https://os-climate.org) - Open source - contributions welcome...



<https://os-climate.org/transition-analysis/>

Physical Risk Modelling Framework



<https://os-climate.org/physical-risk-resilience/>

<https://github.com/os-climate/OS-Climate-Community-Hub>

Add energy usage instrumentation to applications

Carbon peak flattening scheduler

Data lakes to collect energy and carbon measurements

Attribution and allocation algorithms

IoT/Mobile instrumentation and optimization at the edge


Things developers need to build

Energy usage dashboards and reports

Supply chain carbon interchange protocols

Energy to carbon models

Climate change impact and risk models



Questions?



<https://github.com/Green-Software-Foundation/real-time-cloud>

[@adrianco@mastodon.social](https://mastodon.social/@adrianco)

@DevSusOps@hachyderm.io

Cloud Provider Sustainability Current Status and Future Directions

**DevSusOps: Adding sustainability concerns
to development and operations**

Adrian Cockcroft - [OrionX.net](#)
@adrianco@mastodon.social
@DevSusOps@hachyderm.io
October 2023