Microservices Workshop: Why, what, and how to get there

Adrian Cockcroft @adrianco
Technology Fellow - Battery Ventures
October 2016

Battery Ventures

CC BY 4.0 https://creativecommons.org/licenses/by/4.0/legalcode

Workshop vs. Presentation

Questions at any time

Interactive discussions

Share your experiences

Everyone's voice should be heard

What does @adrianco do?

Previously: Netflix, eBay, Sun Microsystems, Cambridge Consultants, City University London - BSc Applied Physics

Topics

Scene Setting
State of the Cloud
What Changes?
Product Processes
Microservices
State of the Art
Segmentation
What's Missing?

Monitoring
Challenges
Migration
Response Times
Serverless
Lock-In
Teraservices
Wrap-Up

Setting the Scene

Why am I here?

By Simon Wardley http://enterpriseitadoption.com/

Why am I here?

By Simon Wardley http://enterpriseitadoption.com/

Why am I here?

By Simon Wardley http://enterpriseitadoption.com/

Why am I here?

By Simon Wardley http://enterpriseitadoption.com/

@adrianco's job at the intersection of cloud and Enterprise IT, looking for disruption and opportunities.

Disruptions in 2016 coming from server-less computing and teraservices.

Typical reactions to my Netflix talks…

Typical reactions to my Netflix talks…

Typical reactions to my Netflix talks…

"You guys are crazy! Can't believe it"
– 2009

"What Netflix is doing won't work"
– 2010

Typical reactions to my Netflix talks…

"You guys are crazy! Can't believe it"
– 2009

"What Netflix is doing won't work"
– 2010

It only works for 'Unicorns' like Netflix"
– 2011

Typical reactions to my Netflix talks…

"You guys are crazy! Can't believe it"
– 2009

"What Netflix is doing won't work"
– 2010

It only works for 'Unicorns' like Netflix"
– 2011

"We'd like to do that but can't"
– 2012

Typical reactions to my Netflix talks…

"You guys are crazy! Can't believe it"
– 2009

"What Netflix is doing won't work"
– 2010

It only works for 'Unicorns' like Netflix"
– 2011

"We'd like to do that but can't"
– 2012

"We're on our way using Netflix OSS code"
– 2013

What I learned from my time at Netflix

What I learned from my time at Netflix

- *Speed wins in the marketplace*

What I learned from my time at Netflix

- *Speed wins in the marketplace*
- *Remove friction from product development*

What I learned from my time at Netflix

- *Speed wins in the marketplace*
- *Remove friction from product development*
- *High trust, low process, no hand-offs between teams*

What I learned from my time at Netflix

- *Speed wins in the marketplace*
- *Remove friction from product development*
- *High trust, low process, no hand-offs between teams*
- *Freedom and responsibility culture*

What I learned from my time at Netflix

- *Speed wins in the marketplace*
- *Remove friction from product development*
- *High trust, low process, no hand-offs between teams*
- *Freedom and responsibility culture*
- *Don’t do your own undifferentiated heavy lifting*

What I learned from my time at Netflix

- *Speed wins in the marketplace*
- *Remove friction from product development*
- *High trust, low process, no hand-offs between teams*
- *Freedom and responsibility culture*
- *Don’t do your own undifferentiated heavy lifting*
- *Use simple patterns automated by tooling*

What I learned from my time at Netflix

- *Speed wins in the marketplace*
- *Remove friction from product development*
- *High trust, low process, no hand-offs between teams*
- *Freedom and responsibility culture*
- *Don’t do your own undifferentiated heavy lifting*
- *Use simple patterns automated by tooling*
- *Self service cloud makes impossible things instant*

“You build it, you run it.”

Werner Vogels 2006

In 2014 Enterprises finally embraced public cloud and in 2015 began replacing entire datacenters.

In 2014 Enterprises finally embraced public cloud and in 2015 began replacing entire datacenters.

Oct 2014

In 2014 Enterprises finally embraced public cloud and in 2015 began replacing entire datacenters.

Oct 2014

Oct 2015

adrian cockcroft @adrianco · Oct 22

RT @devopscouts: Nordstrom went from optimizing for IT cost to optimizing for delivery speed @ladyhock #DevOps #DOES14 < this is key point

20 12

In 2014 Enterprises finally embraced public cloud and in 2015 began replacing entire datacenters.

Oct 2014

Lydia Leong
@cloudpundit

Following

What a difference a year makes. My #GartnerSYM 1:1s this year, everyone's already comfortably using IaaS (overwhelmingly AWS, bit of Azure).

Oct 2015

Lydia Leong @cloudpundit · Oct 7

We're really seeing serious movement of the banks to the cloud at this point. Huge sea change in attitudes.

adrian cockcroft @adrianco

"We can operate more securely on AWS than we can in our own data centers" Rob Alexander of CapitalOne #reinvent

Key Goals of the CIO?
Align IT with the business
Develop products faster
Try not to get breached

Security Blanket Failure

Insecure applications hidden behind firewalls make you feel safe until the breach happens…

http://peanuts.wikia.com/wiki/Linus'_security_blanket

“Web scale” vs. “Enterprise”

"Webscale"
Freedom and responsibility
High trust

“Enterprise”
Bureaucracy and blame
Low trust

How can everyone get speed, low cost, and better usability?

Mixed methods: Disaggregation into microservices helps!

Example Monolith:

Monolithic application

Complex mix of queries

Monolithic "kitchen sink" database

Because one part of the monolithic application and database holds sensitive data all of it is subject to the most rigorous policies

Microservices version:

Segregated team owns secure data sources and infrequent updates

Microservices separation of concerns

Isolated single purpose connections

Optimized datastores

Segregated team owns rapid improvement of most common use cases

Because each microservice can conform to the appropriate policy, demands for agility can be separated from requirements for security

State of the Cloud

Previous Cloud Trend Updates

GigaOM Structure May 2014
D&B Cloud Innovation July 2015
GigaOM Structure November 2015

Previous Cloud Trend Updates

GigaOM Structure May 2014
D&B Cloud Innovation July 2015
GigaOM Structure November 2015

Trends from 2014: Noted as appropriate

Cloud Ecosystem Matures

Staying Power
Support
Scale
Location

Cloud Ecosystem Matures

Staying Power
Support
Scale
Location

Trends from 2014: Even more emphasis on location

Safe Bets

Safe Bets

Safe Bets

Trends from 2014: AWS capacity share vs. Azure still growing
AWS growing 60% year on year, Enterprise share growing fast
Azure providing strong Linux and Open Source support

The Global Land-Grab

Azure	**20** Regions
AWS	**11** Regions
GCE	**6** Regions

http://aws.amazon.com/about-aws/global-infrastructure/
http://azure.microsoft.com/en-us/regions/
https://cloud.google.com/compute/docs/zones#available
http://www.google.com/about/datacenters/inside/locations/index.html

The Global Land-Grab

Azure	**20**	Regions
AWS	**11**	Regions
GCE	**6**	Regions

http://aws.amazon.com/about-aws/global-infrastructure/
http://azure.microsoft.com/en-us/regions/
https://cloud.google.com/compute/docs/zones#available
http://www.google.com/about/datacenters/inside/locations/index.html

The Global Land-Grab

Azure	**20** Regions
AWS	**11** Regions
GCE	**6** Regions

http://aws.amazon.com/about-aws/global-infrastructure/
http://azure.microsoft.com/en-us/regions/
https://cloud.google.com/compute/docs/zones#available
http://www.google.com/about/datacenters/inside/locations/index.html

The Global Land-Grab

Azure	**20** Regions
AWS	**11** Regions
GCE	**6** Regions

http://aws.amazon.com/about-aws/global-infrastructure/
http://azure.microsoft.com/en-us/regions/
https://cloud.google.com/compute/docs/zones#available
http://www.google.com/about/datacenters/inside/locations/index.html

The Global Land-Grab

Azure	**20** Regions
AWS	**11** Regions
GCE	**6** Regions

http://aws.amazon.com/about-aws/global-infrastructure/
http://azure.microsoft.com/en-us/regions/
https://cloud.google.com/compute/docs/zones#available
http://www.google.com/about/datacenters/inside/locations/index.html

The Global Land-Grab

Azure	**20**	Regions
AWS	**11**	Regions
GCE	**6**	Regions

Trends from 2014: AWS and Azure go after big markets
2016 Google finally decides it needs more regions too…

http://aws.amazon.com/about-aws/global-infrastructure/
http://azure.microsoft.com/en-us/regions/
https://cloud.google.com/compute/docs/zones#available
http://www.google.com/about/datacenters/inside/locations/index.html

Who is challenging VMware for in-house cloud automation?

Questions we keep hearing about cloud

Challenges for Private Cloud

Stability
Functionality
Total Cost of Ownership

Challenges for Private Cloud

vmware

openstack™

Stability
Functionality
Total Cost of Ownership

Challenges for Private Cloud

Stability
Functionality
Total Cost of Ownership

Some enterprise vendor responses to cloud and container ecosystem growth…

The ship is sinking, let's re-brand as a submarine!

The ship is sinking, let's re-brand as a submarine!

The ship is sinking, let's merge with a submarine!

EMC² *The ship is sinking, let's merge with a submarine!*

Look! we cut our ship in two really quickly!

Hewlett Packard Enterprise

What needs to change?

Developer responsibilities: Faster, cheaper, safer

“It isn't what we don't know that gives us trouble, it's what we know that ain't so.”

Will Rogers

Assumptions

Optimizations

Assumption: Process prevents problems

Organizations build up slow complex "Scar tissue" processes

"This is the IT swamp draining manual for anyone who is neck deep in alligators."

1984

2014

Product Development Processes

Waterfall Product Development

Waterfall Product Development

Hardware provisioning is undifferentiated heavy lifting – replace it with IaaS

Waterfall Product Development

Hardware provisioning is undifferentiated heavy lifting – replace it with IaaS

Waterfall Product Development

Hardware provisioning is undifferentiated heavy lifting – replace it with IaaS

Process Hand-Off Steps for Agile Development on IaaS

IaaS Agile Product Development

Business Need
- Documents
- Weeks

Software Development
- Specifications
- Weeks

Deployment and Testing
- Reports
- Days

Customer Feedback
- It sucks!
- Days

IaaS Agile Product Development

Business Need
- Documents
- Weeks

Software Development
- Specifications
- Weeks

Deployment and Testing
- Reports
- Days

Customer Feedback
- It sucks!
- Days

IaaS Agile Product Development

etc…

Business Need
- Documents
- Weeks

Software Development
- Specifications
- Weeks

Deployment and Testing
- Reports
- Days

Customer Feedback
- It sucks!
- Days

Software provisioning is undifferentiated heavy lifting – replace it with PaaS

IaaS Agile Product Development

Software provisioning is undifferentiated heavy lifting – replace it with PaaS

IaaS Agile Product Development

etc…

Business Need
- Documents
- Weeks

Software Development
- Specifications
- Weeks

Customer Feedback
- It sucks!
- Days

Software provisioning is undifferentiated heavy lifting – replace it with PaaS

Process for Continuous Delivery of Features on PaaS

PaaS CD Feature Development

Business Need
- Discussions
- Days

Software Development
- Code
- Days

Customer Feedback
- Fix this Bit!
- Hours

PaaS CD Feature Development

Business Need
- Discussions
- Days

Software Development
- Code
- Days

Customer Feedback
- Fix this Bit!
- Hours

Building your own business apps is undifferentiated heavy lifting – use SaaS

PaaS CD Feature Development

Building your own business apps is undifferentiated heavy lifting – use SaaS

PaaS CD Feature Development

etc…

Business Need
- Discussions
- Days

Customer Feedback
- Fix this Bit!
- Hours

Building your own business apps is undifferentiated heavy lifting – use SaaS

SaaS Based Business Application Development

Business Need
- GUI Builder
- Hours

Customer Feedback
- Fix this bit!
- Seconds

SaaS Based Business Application Development

and thousands more…

Business Need
- GUI Builder
- Hours

Customer Feedback
- Fix this bit!
- Seconds

Value Chain Mapping

Simon Wardley http://blog.gardeviance.org/2014/11/how-to-get-to-strategy-in-ten-steps.html
Related tools and training http://www.wardleymaps.com/

Value Chain Mapping

Your unique product - Agile

Simon Wardley http://blog.gardeviance.org/2014/11/how-to-get-to-strategy-in-ten-steps.html
Related tools and training http://www.wardleymaps.com/

Value Chain Mapping

Your unique product - Agile

Best of breed as a Service - Lean

Simon Wardley http://blog.gardeviance.org/2014/11/how-to-get-to-strategy-in-ten-steps.html
Related tools and training http://www.wardleymaps.com/

Value Chain Mapping

Your unique product - Agile

Best of breed as a Service - Lean

Undifferentiated utility suppliers - 6sigma

Simon Wardley http://blog.gardeviance.org/2014/11/how-to-get-to-strategy-in-ten-steps.html
Related tools and training http://www.wardleymaps.com/

Observe
Orient
Decide
Act
Continuous Delivery

Land grab opportunity
Measure Customers
Observe
Competitive Move
Customer Pain Point
Act
Continuous Delivery
Orient
Decide

INNOVATION

Land grab opportunity
Competitive Move
Measure Customers
Customer Pain Point
Observe
Orient
Decide
Act
Continuous Delivery

INNOVATION

Land grab opportunity
Measure Customers
Competitive Move
Customer Pain Point
Observe
Orient
Decide
Act
Continuous Delivery
Analysis
Model Hypotheses

INNOVATION
Land grab opportunity
Measure Customers
Competitive Move
Customer Pain Point
Observe
Orient
Decide
Act
Continuous Delivery
Analysis
BIG DATA
Model Hypotheses

INNOVATION
Land grab opportunity
Competitive Move
Measure Customers
Observe
Customer Pain Point
Act
Continuous Delivery
Orient
Analysis
BIG DATA
Model Hypotheses
Decide
Plan Response
Share Plans
JFDI

INNOVATION
Land grab opportunity
Competitive Move
Measure Customers
Observe
Customer Pain Point
Act
Continuous Delivery
Orient
Analysis
BIG DATA
Model Hypotheses
Decide
Plan Response
Share Plans
JFDI
CULTURE

INNOVATION
Land grab opportunity
Competitive Move
Measure Customers
Observe
Customer Pain Point
Launch AB Test
Automatic Deploy
Act
Continuous Delivery
Orient
Analysis
BIG DATA
Model Hypotheses
Incremental Features
Decide
Plan Response
Share Plans
JFDI
CULTURE

INNOVATION
Land grab opportunity
Competitive Move
Measure Customers
Customer Pain Point
Observe
Launch AB Test
Automatic Deploy
Incremental Features
Act
Continuous Delivery
Orient
Analysis
BIG DATA
Model Hypotheses
CLOUD
Decide
Plan Response
Share Plans
JFDI
CULTURE

INNOVATION
Land grab opportunity
Competitive Move
Measure Customers
Customer Pain Point
Observe
Launch AB Test
Automatic Deploy
Incremental Features
Act
Continuous Delivery
Orient
Analysis
BIG DATA
Model Hypotheses
CLOUD
Decide
Plan Response
Share Plans
JFDI
CULTURE

INNOVATION
Land grab opportunity
Competitive Move
Measure Customers
Customer Pain Point
Observe
Launch AB Test
Automatic Deploy
Incremental Features
Act
Continuous Delivery
Orient
Analysis
BIG DATA
Model Hypotheses
CLOUD
Decide
Plan Response
Share Plans
JFDI
CULTURE

Breaking Down the SILOs

Breaking Down the SILOs

Breaking Down the SILOs

Breaking Down the SILOs

Breaking Down the SILOs

Breaking Down the SILOs

Breaking Down the SILOs

Re-Org from project teams to product teams

Monolithic service updates

Works well with a small number of developers and a single language like php, java or ruby

Monolithic service updates
Release Plan
Developer
Developer
Developer
Developer
Developer
QA Release Integration
Bugs
Ops Replace Old With New Release
Works well with a small number of developers and a single language like php, java or ruby

Monolithic service updates
Release Plan
Developer
Developer
Developer
Developer
Developer
QA Release Integration
Bugs
Ops Replace Old With New Release
Bugs
Works well with a small number of developers and a single language like php, java or ruby

Use monolithic apps for small teams, simple systems and when you must, to optimize for efficiency and latency

Immutable microservice deployment scales, is faster with large teams and diverse platform components

Immutable microservice deployment scales, is faster with large teams and diverse platform components

Immutable microservice deployment scales, is faster with large teams and diverse platform components

Immutable microservice deployment scales, is faster with large teams and diverse platform components

Standardized container deployment saves time and effort

Standardized container deployment saves time and effort

Run What You Wrote

Run What You Wrote

Run What You Wrote

Run What You Wrote

99.95% customer success rate

Run What You Wrote

Run What You Wrote

Non-Destructive Production Updates

- *"Immutable Code" Service Pattern*
  - *Existing services are unchanged, old code remains in service*
  - *New code deploys as a new service group*
  - *No impact to production until traffic routing changes*
- *A|B Tests, Feature Flags and Version Routing control traffic*
  - *First users in the test cell are the developer and test engineers*
  - *A cohort of users is added looking for measurable improvement*

Deliver four features every four weeks

Work In Progress = 4

Opportunity for bugs: 100% (baseline)

Time to debug each: 100% (baseline)

Deliver four features every four weeks

Bugs! Which feature broke?
Need more time to test!
Extend release to six weeks?

Work In Progress = 4
Opportunity for bugs: 100% (baseline)
Time to debug each: 100% (baseline)

Deliver four features every four weeks

Bugs! Which feature broke?
Need more time to test!
Extend release to six weeks?

Work In Progress = 4
Opportunity for bugs: 100% (baseline)
Time to debug each: 100% (baseline)

But: risk of bugs in delivery increases with interactions!

Deliver four features every four weeks

Bugs! Which feature broke?
Need more time to test!
Extend release to six weeks?

Work In Progress = 4
Opportunity for bugs: 100% (baseline)
Time to debug each: 100% (baseline)

But: risk of bugs in delivery increases with interactions!

Deliver six features every six weeks

Deliver six features every six weeks

Work In Progress = 6

Individual bugs: 150%

Interactions: 150%?

Deliver six features every six weeks

More features
What broke?
More interactions
Even more bugs!!

Work In Progress = 6
Individual bugs: 150%
Interactions: 150%?

Deliver six features every six weeks

More features
What broke?
More interactions
Even more bugs!!

Work In Progress = 6
Individual bugs: 150%
Interactions: 150%?

Deliver six features every six weeks

More features
What broke?
More interactions
Even more bugs!!

Work In Progress = 6
Individual bugs: 150%
Interactions: 150%?

Risk of bugs in delivery increased to 225% of original!

Deliver two features every two weeks

Fewer interactions

Fewer bugs

Better flow

Less Work In Progress

Work In Progress = 2

Opportunity for bugs: 50%

Time to debug each: 50%

Complexity of delivery decreased by 75% from original

If it hurts, do it more often!
Change One Thing at a Time!

What Happened?

Low Cost of Change Using Docker

Developers
- Compile/Build
- Seconds

Extend container
- Package dependencies
- Seconds

Deploy Container
- Docker startup
- Seconds

Low Cost of Change Using Docker

Developers
- Compile/Build
- Seconds

Extend container
- Package dependencies
- Seconds

Deploy Container
- Docker startup
- Seconds

Fast tooling supports continuous delivery of many tiny changes

Disruptor: Continuous Delivery with Containerized Microservices

It’s what you know that isn’t so

It’s what you know that isn’t so

- *Make your assumptions explicit*

It’s what you know that isn’t so

- *Make your assumptions explicit*
- *Extrapolate trends to the limit*

It’s what you know that isn’t so

- *Make your assumptions explicit*
- *Extrapolate trends to the limit*
- *Listen to non-customers*

It’s what you know that isn’t so

- *Make your assumptions explicit*
- *Extrapolate trends to the limit*
- *Listen to non-customers*
- *Follow developer adoption, not IT spend*

It’s what you know that isn’t so

- *Make your assumptions explicit*
- *Extrapolate trends to the limit*
- *Listen to non-customers*
- *Follow developer adoption, not IT spend*
- *Map evolution of products to services to utilities*

It’s what you know that isn’t so

- *Make your assumptions explicit*
- *Extrapolate trends to the limit*
- *Listen to non-customers*
- *Follow developer adoption, not IT spend*
- *Map evolution of products to services to utilities*
- *Re-organize your teams for speed of execution*

Nuances of Lock-In

Most IT Ops people will try to avoid lock-in

Most product developers will pick the best of breed option

Analogy:
Dev's get to go on dates and get married
Ops get to run the family and get divorced

DevOps gets the whole lifecycle!

DevOps to the rescue!

https://www.youtube.com/watch?v=7g3uqSzWVZs

Erik Rogneby @erikrogneby · Mar 9

@botchagalupe John, do you mind if we thrown your ignite **Deming** 14 point video up on our DevOps page? Brilliant! youtube.com/watch?v=7g3uqS…

YouTube

John Willis' Ignite talk: "Devops and Dr Deming's 14 Points" (Veloc…

This a rapid paced 20 slide Ignite presentation about Edward Deming and his 14 points. More specifically we will analyze each one of his 14 points as to how …

"End the practice of awarding business on the basis of a price tag. Instead, minimize total cost. Move toward a single supplier for any one item, on a long-term relationship of loyalty and trust."

W. Edwards Deming - 4th Point

"End the practice of awarding business on the basis of a price tag. Instead, minimize total cost. Move toward a single supplier for any one item, on a ~~long-term~~ relationship of ~~loyalty and trust.~~"

dysfunctional *exploitation and abuse*

How did we end up here?

Project vs. Product
Leads to lock-in
Evolves to follow best of breed

Evolution
Technology Refresh
Move to open Source
On-prem -> as a Service

Best of breed is now OSS and as a Service

Less inherent lock-in

What kinds of lock-in are there?

Business lock-in

Hardest to escape…

Legal lock-in

e.g. compliance with laws that exclude alternatives based on jurisdiction or certification

Legal lock-in

e.g. compliance with laws that exclude alternatives based on jurisdiction or certification

Financial lock-in

e.g. budget spent in advance on long term deal with a vendor

Legal lock-in

e.g. compliance with laws that exclude alternatives based on jurisdiction or certification

Financial lock-in

e.g. budget spent in advance on long term deal with a vendor

Contractual lock-in

e.g. partnership or investment deal with one vendor prevents using alternatives

Technology lock-in

Possible to escape given time and work…

Proximity lock-in

e.g. chatty clients don't work unless they are co-located with their server

Proximity lock-in

e.g. chatty clients don't work unless they are co-located with their server

Topology lock-in

e.g. quorum based availability (C, Riak) needs three zones/datacenters per region*

Implementation

e.g. interface is the same but behavior is different

Proximity lock-in

e.g. chatty clients don't work unless they are co-located with their server

Topology lock-in

e.g. quorum based availability (C, Riak) needs three zones/datacenters per region*

Soft lock-in

Relatively easy to escape…

Interface lock-in

e.g. different APIs that get the same result, easy to hide behind an abstraction layer

Interface lock-in

e.g. different APIs that get the same result, easy to hide behind an abstraction layer

Query syntax lock-in

e.g. SQL variants for different databases

Interface lock-in

e.g. different APIs that get the same result, easy to hide behind an abstraction layer

Web service lock-in

Interface lock-in, but remote access unlocks ability to migrate applications

Query syntax lock-in

e.g. SQL variants for different databases

Interface lock-in

e.g. different APIs that get the same result, easy to hide behind an abstraction layer

Web service lock-in

Interface lock-in, but remote access unlocks ability to migrate applications

Query syntax lock-in

e.g. SQL variants for different databases

Data gravity lock-in

e.g. lots of data to move or duplicate

Cloud native microservices

Example for discussion

Example for discussion

Example for discussion

Example for discussion

Example for discussion

Example for discussion

Example for discussion

Microservices

A Microservice Definition

Loosely coupled service oriented architecture with bounded contexts

If every service has to be updated at the same time it's not loosely coupled

A Microservice Definition

Loosely coupled service oriented architecture with bounded contexts

If every service has to be updated at the same time it's not loosely coupled

A Microservice Definition

Loosely coupled service oriented architecture with bounded contexts

If you have to know too much about surrounding services you don't have a bounded context. See the Domain Driven Design book by Eric Evans.

Coupling Concerns

- *Conway's Law - organizational coupling*
- *Centralized Database Schemas*
- *Enterprise Service Bus - centralized message queues*
- *Inflexible Protocol Versioning*

http://en.wikipedia.org/wiki/Conway's_law

Speeding Up The Platform

Datacenter Snowflakes

- Deploy in months
- Live for years

Speeding Up The Platform

Datacenter Snowflakes
- Deploy in months
- Live for years

Virtualized and Cloud
- Deploy in minutes
- Live for weeks

Speeding Up The Platform

Datacenter Snowflakes
- Deploy in months
- Live for years

Virtualized and Cloud
- Deploy in minutes
- Live for weeks

Container Deployments
- Deploy in seconds
- Live for minutes/hours

Speeding Up The Platform

Datacenter Snowflakes
- Deploy in months
- Live for years

Virtualized and Cloud
- Deploy in minutes
- Live for weeks

Container Deployments
- Deploy in seconds
- Live for minutes/hours

Lambda Deployments
- Deploy in milliseconds
- Live for seconds

Speeding Up The Platform

Datacenter Snowflakes
- Deploy in months
- Live for years

Virtualized and Cloud
- Deploy in minutes
- Live for weeks

Container Deployments
- Deploy in seconds
- Live for minutes/hours

Lambda Deployments
- Deploy in milliseconds
- Live for seconds

AWS Lambda is leading exploration of serverless architectures in 2016

Separate Concerns with Microservices

- *Invert Conway's Law – teams own service groups and backend stores*
- *One "verb" per single function micro-service, size doesn't matter*
- *One developer independently produces a micro-service*
- *Each micro-service is it's own build, avoids trunk conflicts*
- *Deploy in a container: Tomcat, AMI or Docker, whatever…*
- *Stateless business logic. Cattle, not pets.*
- *Stateful cached data access layer using replicated ephemeral instances*

http://en.wikipedia.org/wiki/Conway's_law

Inspiration

State of the Art in Web Scale Microservice Architectures

AWS Re:Invent : Asgard to Zuul https://www.youtube.com/watch?v=p7ysHhs5hl0
Resiliency at Massive Scale https://www.youtube.com/watch?v=ZfYJHtVL1_w
Microservice Architecture https://www.youtube.com/watch?v=CriDUYtfrjs
New projects for 2015 and Docker Packaging https://www.youtube.com/watch?v=hi7BDAtjfKY
Spinnaker deployment pipeline https://www.youtube.com/watch?v=dwdVwE52KkU
http://www.infoq.com/presentations/spring-cloud-2015

http://www.infoq.com/presentations/Twitter-Timeline-Scalability
http://www.infoq.com/presentations/twitter-soa
http://www.infoq.com/presentations/Zipkin
http://www.infoq.com/presentations/scale-gilt

Go-Kit https://www.youtube.com/watch?v=aL6sd4d4hxk
http://www.infoq.com/presentations/circuit-breaking-distributed-systems
https://speakerdeck.com/mattheath/scaling-micro-services-in-go-highload-plus-plus-2014

Microservice Architectures

NETFLIX | OSS Microservices

Spinnaker SpringCloud Tooling	Edda Archaius Configuration	Eureka Prana Discovery	Denominator Zuul Ribbon Routing	Hystrix Pytheus Atlas Observability

Ephemeral datastores using Dynomite, Memcached, Astyanax, Staash, Priam, Cassandra

Orchestration with Autoscalers on AWS, Titus exploring Mesos & ECS for Docker

Development using Java, Groovy, Scala, Clojure, Python with AMI and Docker Containers

Policy via the Simian Army - Chaos Monkey, Chaos Gorilla, Conformity Monkey, Security Monkey

Cloud Native Storage

Cloud Native Storage

Cloud Native Storage

SSDs inside arrays disrupt incumbent suppliers

Cloud Native Storage

SSDs inside arrays disrupt incumbent suppliers

SSDs inside ephemeral instances disrupt an entire industry

Cloud Native Storage

SSDs inside arrays disrupt incumbent suppliers

SSDs inside ephemeral instances disrupt an entire industry

NetflixOSS Uses Priam to create Cassandra clusters in minutes

NETFLIX | OSS Netflix Open Source Software Center

Repositories | Powered By NetflixOSS

These companies are using and contributing to Netflix OSS Components

Email netflixoss@netflix.com to have your logo here.

NETFLIX

Twitter Microservices

Tooling	Decider Configuration	Finagle Zookeeper Discovery	Finagle Netty Routing	Zipkin Observability

Custom Cassandra-like datastore: Manhattan

Orchestration using Aurora deployment in datacenters using Mesos

Scala with JVM Container

Twitter Microservices

Tooling

Decider

Configuration

Finagle
Zookeeper

Discovery

Finagle
Netty

Routing

Zipkin

Observability

Custom Cassandra-like datastore: Manhattan

Orchestration using Aurora deployment in datacenters using Mesos

Scala with JVM Container

Focus on efficient datacenter deployment at scale

Gilt Microservices

Ion Cannon
SBT
Rake

Tooling

Decider

Configuration

Finagle
Zookeeper

Discovery

Akka
Finagle
Netty

Routing

Zipkin

Observability

Datastores per Microservice using MongoDB, Postgres, Voldemort

Deployment on AWS

Scala and Ruby with Docker Containers

Gilt Microservices

Ion Cannon SBT Rake Tooling	Decider Configuration	Finagle Zookeeper Discovery	Akka Finagle Netty Routing	Zipkin Observability

Datastores per Microservice using MongoDB, Postgres, Voldemort

Deployment on AWS

Scala and Ruby with Docker Containers

Focus on fast development with Scala and Docker

GILT

Hailo Microservices

Hubot
Janky
Jenkins

Tooling

Configuration

go-platform

Discovery

go-platform
RabbitMQ

Routing

Observability

Deployment on AWS

Deployment on AWS

Go using AMI Container and Docker

Hailo Microservices

See: go-micro and https://github.com/peterbourgon/gokit

hailo-2-api
service.login
infrastructure.compute
infrastructure.whisper
kernel.binding
kernel.discovery
kernel.provisioning
service.account
service.adhoc-payment
service.allocation
service.app-feedback
service.app-version
service.archived-job
service.b2c-events
service.b2c-profile
service.banning
service.beta-customer
service.billing-shim
service.card-cancellation
service.card-gateway
service.charging
service.city
service.comms
service.config
service.control-updater
service.corporate-account
service.crash
service.crm-preference
service.customer
service.dispatcher
service.dispatcher-other
service.dispatch-job
service.driver
service.driver-assignment
service.driver-heatmap
service.driver-location
service.driver-organization
service.driver-scoring
service.eta
service.event-notifier
service.experiment
service.fare
service.failure-flag
service.feature-flag
service.hub
service.hub-b2b-rm
service.hub-organization
service.hub-reporting
service.hub-transaction
service.hub-user
service.hub-utility
service.hub-editor
service.hub-event
service.hub-location
service.hub-login
service.job
service.job-events
service.job-gaiota
service.job-system-cancel
service.log
service.log
service.nearest-driver
service.note
service.outage-impact
service.organization
service.payment-validation
service.phone
service.pricing
service.profile
service.profile-flow
service.provisioning-manager
service.public-route
service.push
service.rapid
service.rapid-impact
service.rapid-quote
service.rapid-rate
service.rapid-rate-history
service.rapid-rate-history-v2
service.rapid-rate-v2
service.rapid-rate-v3
service.rapid-rate-v4
service.rapid-rate-v5
service.rapid-rate-v6
service.rapid-rate-v7
service.rapid-rate-v8
service.rapid-rate-v9
service.rapid-rate-v10
service.rapid-rate-v11
service.rapid-rate-v12
service.rapid-rate-v13
service.rapid-rate-v14
service.rapid-rate-v15
service.rapid-rate-v16
service.rapid-rate-v17
service.rapid-rate-v18
service.rapid-rate-v19
service.rapid-rate-v20
service.rapid-rate-v21
service.rapid-rate-v22
service.rapid-rate-v23
service.rapid-rate-v24
service.rapid-rate-v25
service.rapid-rate-v26
service.rapid-rate-v27
service.rapid-rate-v28
service.rapid-rate-v29
service.rapid-rate-v30
service.rapid-rate-v31
service.rapid-rate-v32
service.rapid-rate-v33
service.rapid-rate-v34
service.rapid-rate-v35
service.rapid-rate-v36
service.rapid-rate-v37
service.rapid-rate-v38
service.rapid-rate-v39
service.rapid-rate-v40
service.rapid-rate-v41

Next Generation Applications

Fill in the gaps, rapidly evolving ecosystem choices

Docker
Desktop
Spinnaker
Terraform

Tooling

Wasabi
LaunchDarkly
Habitat

Configuration

Etcd
Eureka
Consul

Discovery

Compose
Linkerd
Weave

Routing

OpenTracing
Prometheus
Hystrix

Observability

Datastores: Orchestrated, Distributed Ephemeral e.g. Cassandra, or DBaaS e.g. DynamoDB

Operational: Mesos, Kubernetes, Swarm, Nomad for private clouds. ECS, GKS for public

Development: Containers and languages e.g. Docker, Lambda, Node.js, Go, Rust

Policy: Security compliance e.g. Vault, Docker Content Trust. Scaffolding e.g. Cloud Foundry

Segmentation

Industry Trends

Airgaps closing
Industrial IoT
Security blanket perimeter firewalls
Datacenter to cloud transitions
New systems of engagement

http://peanuts.wikia.com/wiki/Linus'_security_blanket

Policy

A can talk to B
B can talk to C
A must not talk to C

Y and Z failure modes must be independent so X can always succeed

Availability requirements drive a need for distributed segmentation

Choices?

Too many
choices!

Over-reliance on one mechanism leads to abuse...

Lack of coordination across many mechanisms leads to fragility

Example segmentation mechanisms

Accounts and Roles

Who can set policy for what?
Needs distributed policy management

Network Segmentation

Who controls the network?

Network Segmentation

Datacenter policies are based on separation of duties. Tickets, Network admins and VLANs

Network Segmentation

AWS VPC networking uses developer-driven automation, loses separation of duties…

VPC Abuse Antipattern

Lots of small VPC networks for microservices, end up in IP address space capacity management hell...

Hypervisor and Security Group Segmentation

Distributed firewall rules

Security Group Abuse Antipattern

Too many microservices need to be in the same group, overloads configuration limitations

Kernel eBPF & Calico IPtables Segmentation

Distributed firewall rules

IPtables Segmentation

Can use IP Sets to scale
Managed in the container host OS
Separates routing reachability from access policy

Docker & Weave Segmentation

Docker daemon manages connections

Docker Compose V1

```
proxy:
  build: ./proxy
  ports:
    - "8080:8080"
  links:
    - app
app:
  build: ./app
  links:
    - db

db:
  image: postgres
```

Docker Compose V2

```yaml
version: '2'

services:
  proxy:
    build: ./proxy
    ports:
      - "8080:8080"
    networks:
      - front
  app:
    build: ./app
    networks:
      - front
      - back
  db:
    image: postgres
    networks:
      - back

networks:
  front:
  back:
```

Docker Segmentation

Overlay network created and managed by Docker or Weave. DNS based lookups.

Segmentation Scalability
Real world microservices architectures have hundreds to thousands of distinct microservices

Segmentation Scalability
There's often a few very popular microservices that everyone else wants to talk to

In Search of Segmentation

Datacenters
AD/LDAP Roles
VLAN Networks
Hypervisor
IPtables
Docker Links

In Search of Segmentation

Datacenters	AWS Accounts
AD/LDAP Roles	IAM Roles
VLAN Networks	VPC
Hypervisor	Security Groups
IPtables	Calico Policy
Docker Links	Docker Net/Weave

Hierarchical Segmentation

Setup all the layers with Terraform?

What's Missing?

Advanced Microservices Topics

Failure injection testing
Versioning, routing
Binary protocols and interfaces
Timeouts and retries
Denormalized data models
Monitoring, tracing
Simplicity through symmetry

Failure Injection Testing

Netflix Chaos Monkey, Simian Army, FIT and Gremlin

http://techblog.netflix.com/2011/07/netflix-simian-army.html
http://techblog.netflix.com/2014/10/fit-failure-injection-testing.html
http://techblog.netflix.com/2016/01/automated-failure-testing.html
https://www.infoq.com/presentations/failure-test-research-netflix

@adrianco
BV
Battery Ventures

Trust with Verification

- *Chaos Monkey - enforcing stateless business logic*
- *Chaos Gorilla - enforcing zone isolation/replication*
- *Chaos Kong - enforcing region isolation/replication*
- *Security Monkey - watching for insecure configuration settings*
- *FIT & Gremlin - inject errors to enforce robust dependencies*
- *See over 100 NetflixOSS projects at netflix.github.com*
- *Get "Technical Indigestion" reading techblog.netflix.com*

Benefits of version aware routing

Immediately and safely introduce a new version
Canary test in production
Use DIY feature flags, LaunchDarkly, A|B tests with Wasabi

Route clients to a version so they can't get disrupted
Change client or dependencies but not both at once

Eventually remove old versions
Incremental or infrequent "break the build" garbage collection

Versioning, Routing

Version numbering: Interface.Feature.Bugfix

V1.2.3 to V1.2.4 - Canary test then remove old version

V1.2.x to V1.3.x - Canary test then remove or keep both
Route V1.3.x clients to new version to get new feature
Remove V1.2.x only after V1.3.x is found to work for V1.2.x clients

V1.x.x to V2.x.x - Route clients to specific versions
Remove old server version when all old clients are gone

Protocols

Measure serialization, transmission, deserialization costs

Sending a megabyte of XML between microservices will make you sad, but not as sad as 10yrs ago with SOAP

Use Thrift, Protobuf/gRPC, Avro, SBE internally
Use JSON for external/public interfaces
https://github.com/real-logic/simple-binary-encoding

Interfaces

When you build a service, build a "driver" client for it
Reference implementation error handling and serialization
Release automation stress test using client
Validate that service interface is usable!
Minimize additional dependencies

Swagger - OpenAPI Specification
Datawire Quark adds behaviors to API spec

Interface Version Pinning

Change one thing at a time!
Pin the version of everything else

Incremental build/test/deploy pipeline

Deploy existing app code with new platform
Deploy existing app code with new dependencies
Deploy new app code with pinned platform/dependencies

Interfaces between teams

Interfaces between teams

Interfaces between teams

Interfaces between teams

Interfaces between teams

Interfaces between teams

Interfaces between teams

Interfaces between teams

Interfaces between teams

Interfaces between teams

@adrianco
BV
Battery Ventures

Timeouts and Retries

Connection timeout vs. request timeout confusion

Usually setup incorrectly, global defaults

Systems collapse with "retry storms"

Timeouts too long, too many retries

Services doing work that can never be used

Connections and Requests

TCP makes a connection, HTTP makes a request
HTTP hopefully reuses connections for several requests

Both have different timeout and retry needs!

TCP timeout is purely a property of one network latency hop
HTTP timeout depends on the service and its dependencies

Timeouts and Retries

Bad config: Every service defaults to 2 second timeout, two retries

If anything breaks, everything upstream stops responding

Retries add unproductive work

Timeouts and Retries

Bad config: Every service defaults to 2 second timeout, two retries

If anything breaks, everything upstream stops responding

Retries add unproductive work

Timeouts and Retries

Bad config: Every service defaults to 2 second timeout, two retries

If anything breaks, everything upstream stops responding

Retries add unproductive work

Timeouts and Retries

Bad config: Every service defaults to 2 second timeout, two retries

Timeouts and Retries

Bad config: Every service defaults to 2 second timeout, two retries

First request from Edge timed out so it ignores the successful response and keeps retrying. Middle service load increases as it's doing work that isn't being consumed

Timeouts and Retries

Bad config: Every service defaults to 2 second timeout, two retries

First request from Edge timed out so it ignores the successful response and keeps retrying. Middle service load increases as it's doing work that isn't being consumed

Timeout and Retry Fixes

Cascading timeout budget
Static settings that decrease from the edge
or dynamic budget passed with request

How often do retries actually succeed?
Don't ask the same instance the same thing
Only retry on a different connection

Timeouts and Retries

Budgeted timeout, one retry

Timeouts and Retries

Budgeted timeout, one retry

*Upstream timeout must always be longer than total downstream timeout * retries delay*

No unproductive work while fast failing

Timeouts and Retries

Budgeted timeout, failover retry

For replicated services with multiple instances never retry against a failed instance

No extra retries or unproductive work

Timeouts and Retries

Budgeted timeout, failover retry

For replicated services with multiple instances never retry against a failed instance

No extra retries or unproductive work

Manage Inconsistency

ACM Paper: "The Network is Reliable"
Distributed systems are inconsistent by nature
Clients are inconsistent with servers
Most caches are inconsistent
Versions are inconsistent
Get over it and
Deal with it

See http://queue.acm.org/detail.cfm?id=2655736

Denormalized Data Models

Any non-trivial organization has many databases
Cross references exist, inconsistencies exist

Microservices work best with individual simple stores
Scale, operate, mutate, fail them independently

NoSQL allows flexible schema/object versions

Denormalized Data Models

Build custom cross-datasource check/repair processes
Ensure all cross references are up to date
Read these Pat Helland papers

Immutability Changes Everything
http://highscalability.com/blog/2015/1/26/paper-immutability-changes-everything-by-pat-helland.html

Memories, Guesses and Apologies
https://blogs.msdn.microsoft.com/pathelland/2007/05/15/memories-guesses-and-apologies/

Standing on the Distributed Shoulders of Giants
http://queue.acm.org/detail.cfm?id=2953944

Cloud Native Monitoring and Microservices

Cloud Native Microservices

- *High rate of change*
  - *Code pushes can cause floods of new instances and metrics*
  - *Short baseline for alert threshold analysis – everything looks unusual*
- *Ephemeral Configurations*
  - *Short lifetimes make it hard to aggregate historical views*
  - *Hand tweaked monitoring tools take too much work to keep running*
- *Microservices with complex calling patterns*
  - *End-to-end request flow measurements are very important*
  - *Request flow visualizations get overwhelmed*

Microservice Based Architectures

See http://www.slideshare.net/LappleApple/gilt-from-monolith-ruby-app-to-micro-service-scala-service-architecture

Continuous Delivery and DevOps

- *Changes are smaller but more frequent*
- *Individual changes are more likely to be broken*
- *Changes are normally deployed by developers*
- *Feature flags are used to enable new code*
- *Instant detection and rollback matters much more*

Whoops! I didn't mean that! Reverting…

Not cool if it takes 5 minutes to see it failed and 5 more to see a fix
No-one notices if it only takes 5 seconds to detect and 5 to see a fix

NetflixOSS Hystrix/Turbine Circuit Breaker

http://techblog.netflix.com/2012/12/hystrix-dashboard-and-turbine.html

NetflixOSS Hystrix/Turbine Circuit Breaker

http://techblog.netflix.com/2012/12/hystrix-dashboard-and-turbine.html

Low Latency SaaS Based Monitors

https://www.datadoghq.com/ http://www.instana.com/ www.bigpanda.io www.vividcortex.com signalfx.com wavefront.com sysdig.com dataloop.io

See www.battery.com for a list of portfolio investments

A Tragic Quadrant

A Tragic Quadrant

YMMV: Opinionated approximate positioning only

Metric to display latency needs to be less than human attention span (~10s)

Challenges for Microservice Platforms

Managing Scale

It's much more challenging than just a large number of machines

A Possible Hierarchy	***How Many?***
Continents	*3 to 5*
Regions	*2-4 per Continent*
Zones	*1-5 per Region*
Services	*100's per Zone*
Versions	*Many per Service*
Containers	*1000's per Version*
Instances	*10,000's*

Flow

Some tools can show the request flow across a few services

APPDYNAMICS

NETFLIX

Interesting architectures have a lot of microservices! Flow visualization is a big challenge.

See http://www.slideshare.net/LappleApple/gilt-from-monolith-ruby-app-to-micro-service-scala-service-architecture

Simulated Microservices

Model and visualize microservices
Simulate interesting architectures
Generate large scale configurations
Eventually stress test real tools

Code: github.com/adrianco/spigo
Simulate Protocol Interactions in Go
Visualize with D3
See for yourself: http://simianviz.surge.sh
Follow @simianviz for updates

Spigo Nanoservice Structure

```
func Start(listener chan gotocol.Message) {
        ...
        for {
                select {
                case msg := <-listener:
                        flow.Instrument(msg, name, hist)
                        switch msg.Imposition {
                        case gotocol.Hello:      // get named by parent
                                ...
                        case gotocol.NameDrop:   // someone new to talk to
                                ...
                        case gotocol.Put: // upstream request handler
                                ...
                                outmsg := gotocol.Message{gotocol.Replicate, listener, time.Now(),
                                  msg.Ctx.NewParent(), msg.Intention}
                                flow.AnnotateSend(outmsg, name)
                                outmsg.GoSend(replicas)
                        }
                case <-eurekaTicker.C:           // poll the service re
                        ...
                }
        }
}
```

Instrument incoming requests

update trace context

Instrument outgoing requests

Skeleton code for replicating a Put message

Flow Trace Records

```
us-east-1.zoneC.riak2  t98p895s896 Put
us-east-1.zoneA.riak3  t98p896s908 Replicate
us-east-1.zoneB.riak4  t98p896s909 Replicate
us-west-2.zoneA.riak9  t98p896s910 Replicate
us-west-2.zoneB.riak10 t98p910s912 Replicate
us-west-2.zoneC.riak8  t98p910s913 Replicate
```

Open Zipkin

A common format for trace annotations
A Java tool for visualizing traces
Standardization effort to fold in other formats
Driven by Adrian Cole (currently at Pivotal)
Extended to load Spigo generated trace files

Zipkin Trace Dependencies

Zipkin Trace Dependencies

Trace for one Spigo Flow

Global Visualization

NetflixOSS
Vizceral

https://github.com/Netflix/vizceral
https://github.com/adrianco/go-vizceral

Regional Visualization

NetflixOSS
Vizceral

Service Context Visualization

NetflixOSS
Vizceral

Simple Architecture Principles

Symmetry
Invariants
Stable assertions
No special cases

Migrating to Microservices

See for yourself: http://simianviz.surge.sh/migration

Migrating to Microservices

See for yourself: http://simianviz.surge.sh/migration

Step 1 - Add Memcache

Step 2 - Add Web Proxy Service

Migrating to Microservices

See for yourself: http://simianviz.surge.sh/migration

Step 3 - Add Data Access Layer

Step 4 - Add Microservices

Migrating to Microservices

See for yourself: http://simianviz.surge.sh/migration

Step 5 - Add Cassandra

Step 6 - Remove MySQL

Migrating to Microservices

See for yourself: http://simianviz.surge.sh/migration

Step 7 - Add Second Region

Step 8 - Connect Cassandra Regions

Migrating to Microservices

See for yourself: http://simianviz.surge.sh/migration

Step 9 - Add Third Region

Definition of an architecture

Header includes chaos monkey victim

See for yourself: http://simianviz.surge.sh/lamp

```
{
    "arch": "lamp",
    "description":"Simple LAMP stack",
    "version": "arch-0.0",
    "victim": "webserver",
    "services": [
        { "name": "rds-mysql",     "package": "store",       "count": 2,  "regions": 1, "dependencies": [] },
        { "name": "memcache",      "package": "store",       "count": 1,  "regions": 1, "dependencies": [] },
        { "name": "webserver",     "package": "monolith",    "count": 18, "regions": 1, "dependencies": ["memcache", "rds-mysql"] },
        { "name": "webserver-elb", "package": "elb",         "count": 0,  "regions": 1, "dependencies": ["webserver"] },
        { "name": "www",           "package": "denominator", "count": 0,  "regions": 0, "dependencies": ["webserver-elb"] }
    ]
}
```

New tier name

Tier package

Node count

0 = non Regional

List of tier dependencies

Running Spigo

-a architecture lamp
-j graph json/lamp.json
-d run for 2 seconds

```
$ ./spigo -a lamp -j -d 2
2016/01/26 23:04:05 Loading architecture from json_arch/lamp_arch.json
2016/01/26 23:04:05 lamp.edda: starting
2016/01/26 23:04:05 Architecture: lamp Simple LAMP stack
2016/01/26 23:04:05 architecture: scaling to 100%
2016/01/26 23:04:05 lamp.us-east-1.zoneB.eureka01....eureka.eureka: starting
2016/01/26 23:04:05 lamp.us-east-1.zoneA.eureka00....eureka.eureka: starting
2016/01/26 23:04:05 lamp.us-east-1.zoneC.eureka02....eureka.eureka: starting
2016/01/26 23:04:05 Starting: {rds-mysql     store 1 2 []}
2016/01/26 23:04:05 Starting: {memcache     store 1 1 []}
2016/01/26 23:04:05 Starting: {webserver     monolith 1 18 [memcache rds-mysql]}
2016/01/26 23:04:05 Starting: {webserver-elb     elb 1 0 [webserver]}
2016/01/26 23:04:05 Starting: {www     denominator 0 0 [webserver-elb]}
2016/01/26 23:04:05 lamp.*.*.www00....www.denominator activity rate  10ms
2016/01/26 23:04:06 chaosmonkey delete: lamp.us-east-1.zoneC.webserver02....webserver.monolith
2016/01/26 23:04:07 asgard: Shutdown
2016/01/26 23:04:07 lamp.us-east-1.zoneB.eureka01....eureka.eureka: closing
2016/01/26 23:04:07 lamp.us-east-1.zoneA.eureka00....eureka.eureka: closing
2016/01/26 23:04:07 lamp.us-east-1.zoneC.eureka02....eureka.eureka: closing
2016/01/26 23:04:07 spigo: complete
2016/01/26 23:04:07 lamp.edda: closing
```

Riak IoT Architecture


```
{
    "arch": "riak",
    "description":"Riak IoT ingestion example for the RICON 2015 presentation",
    "version": "arch-0.0",
    "victim": "",
    "services": [
        { "name": "riakTS",      "package": "riak",           "count": 6, "regions": 1, "dependencies": ["riakTS", "eureka"]},
        { "name": "ingester",    "package": "staash",         "count": 6, "regions": 1, "dependencies": ["riakTS"]},
        { "name": "ingestMQ",    "package": "karyon",         "count": 3, "regions": 1, "dependencies": ["ingester"]},
        { "name": "riakKV",      "package": "riak",           "count": 3, "regions": 1, "dependencies": ["riakKV"]},
        { "name": "enricher",    "package": "staash",         "count": 6, "regions": 1, "dependencies": ["riakKV", "ingestMQ"]},
        { "name": "enrichMQ",    "package": "karyon",         "count": 3, "regions": 1, "dependencies": ["enricher"]},
        { "name": "analytics",   "package": "karyon",         "count": 6, "regions": 1, "dependencies": ["ingester"]},
        { "name": "analytics-elb", "package": "elb",          "count": 0, "regions": 1, "dependencies": ["analytics"]},
        { "name": "analytics-api", "package": "denominator", "count": 0, "regions": 0, "dependencies": ["analytics-elb"]},
        { "name": "normalization", "package": "karyon",       "count": 6, "regions": 1, "dependencies": ["enrichMQ"]},
        { "name": "iot-elb",     "package": "elb",            "count": 0, "regions": 1, "dependencies": ["normalization"]},
        { "name": "iot-api",     "package": "denominator",    "count": 0, "regions": 0, "dependencies": ["iot-elb"]},
        { "name": "stream",      "package": "karyon",         "count": 6, "regions": 1, "dependencies": ["ingestMQ"]},
        { "name": "stream-elb", "package": "elb",             "count": 0, "regions": 1, "dependencies": ["stream"]},
        { "name": "stream-api", "package": "denominator",     "count": 0, "regions": 0, "dependencies": ["stream-elb"]}
    ]
}
```


New tier name

Tier package

Node count

0 = non Regional

List of tier dependencies

Single Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

Single Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

Single Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

Single Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

Single Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

Single Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

Single Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

Two Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

Two Region Riak IoT

See for yourself: http://simianviz.surge.sh/riak

What's the response time of the stream endpoint?

Response Times

What's the response time distribution of a very simple storage backed web service?

See http://www.getguesstimate.com/models/1307

https://www.getguesstimate.com/models/1307

Guesstimate

Plans

Storage Service Response Time

View Options

Model of a storage service showing how cache hits and misses contribute to a multi-modal response time distribution. The positions of the modes depend on the relative response times of each path through the model. The amplitude of each mode depends on the cache hit rates through the system. Mean/Median/Percentiles are a poor and unstable characterization of the distribution.

Created by Adrian Cockcroft - @adrianco

Over-all response
580 ± 480
mean 581.70
stdev 332.98
Samples 5000
=GR+((randoml

Staash CPU
72 ± 24
mean 72.10
stdev 14.93
Samples 5000

Memcache Response
220 ± 18
mean 223.56
stdev 10.87
Samples 5000

Memcache Network
160 ± 10
mean 159.76
stdev 6.13
Samples 5000

Memcache CPU
64 ± 15
mean 63.85
stdev 9.17
Samples 5000

Memcache Hit %
40

MySQL hit %
70

MySQL Response
700 ± 440
mean 702.20
stdev 305.07
Samples 5000

MySQL Network
160 ± 10
mean 159.88
stdev 6.15
Samples 5000

MySQL CPU
250 ± 50
mean 247.86
stdev 30.76
Samples 5000

MySQL Disk
660 ± 450
mean 659.52
stdev 292.25
Samples 5000

Hit rates: memcached 40% mysql 70%

Over-all response

memcached hit mode

580 ± 480

mysql cache hit mode

mysql disk access mode

200 400 600 800 1K 1.2K 1.4K 1.6K 1.8K 2K 2.2K 2.4K 2.6K 2.8K

Percentiles

1%	262.469
5%	274.988
50%	481.226
95%	1226.717
99%	1617.009

service cpu time

memcached hit %

=GR+((randomInt(0,100)>IY)?QW:LR)

memcached response

mysql response

Reasoning

CPU time plus memcache hit lookup time is the first mode, second mode is memcache miss MySQL lookup with in-memory result, and third mode is MySQL lookup including disk access. Change the memcached and mysql hit rate %age to see how the height of each mode varies.

Hit rates: memcached 60% mysql 70%

Over-all response

480 ± 410

Percentiles

1%	261.411
5%	272.255
50%	311.920
95%	1098.009
99%	1458.043

=GR+((randomInt(0,100)>IY)?QW:LR)

$f(x)$

Reasoning

CPU time plus memcache hit lookup time is the first mode, second mode is memcache miss MySQL lookup with in-memory result, and third mode is MySQL lookup including disk access. Change the memcached and mysql hit rate %age to see how the height of each mode varies.

Hit rates: memcached 20% mysql 90%

=GR+((randomInt(0,100)>IY)?QW:LR)

Reasoning

CPU time plus memcache hit lookup time is the first mode, second mode is memcache miss MySQL lookup with in-memory result, and third mode is MySQL lookup including disk access. Change the memcached and mysql hit rate %age to see how the height of each mode varies.

Measuring Response Time With Histograms

Changes made to codahale/hdrhistogram

Changes made to go-kit/kit/metrics

Implementation in adrianco/spigo/collect

What to measure?

What to measure?

Spigo Histogram Results

Collected with: `% spigo -d 60 -j -a storage -c`

```
name: storage.*.*..load00...load.denominator_serv
quantiles: [{50 47103} {99 139263}]
From     To       Count  Prob     Bar
20480    21503    2      0.0007  :
21504    22527    2      0.0007  |
23552    24575    1      0.0003  :
24576    25599    5      0.0017  |
25600    26623    5      0.0017  |
26624    27647    1      0.0003  |
27648    28671    3      0.0010  |
28672    29695    5      0.0017  |
29696    30719    127    0.0421  |####
30720    31743    126    0.0418  |####
31744    32767    74     0.0246  |##
32768    34815    281    0.0932  |#########
34816    36863    201    0.0667  |######
36864    38911    156    0.0518  |#####
38912    40959    185    0.0614  |######
40960    43007    147    0.0488  |####
43008    45055    161    0.0534  |#####
45056    47103    125    0.0415  |####
47104    49151    135    0.0448  |####
49152    51199    99     0.0328  |###
51200    53247    82     0.0272  |##
53248    55295    77     0.0255  |##
55296    57343    66     0.0219  |##
57344    59391    54     0.0179  |#
59392    61439    37     0.0123  |#
61440    63487    45     0.0149  |#
63488    65535    33     0.0109  |#
65536    69631    63     0.0209  |##
69632    73727    98     0.0325  |###
73728    77823    92     0.0305  |###
77824    81919    112    0.0372  |###
81920    86015    88     0.0292  |##
86016    90111    55     0.0182  |#
90112    94207    38     0.0126  |#
94208    98303    51     0.0169  |#
98304    102399   32     0.0106  |#
102400   106495   35     0.0116  |#
106496   110591   17     0.0056  |
110592   114687   19     0.0063  |
114688   118783   18     0.0060  |
118784   122879   6      0.0020  |
122880   126975   8      0.0027  |
```

Median and 99th percentile values

service time for load generator

Response time distribution measured in nanoseconds using High Dynamic Range Histogram

Normalized probability

:# Zero counts skipped
|# Contiguous buckets

Go-Kit Histogram Example

Nanoseconds resolution!

```
const (
	maxHistObservable = 1000000 // one millisecond
	sampleCount       = 1000    // data points will be sampled 5000 times to build a distribution by guesstimate
)

var sampleMap map[metrics.Histogram][]int64
var sampleLock sync.Mutex

func NewHist(name string) metrics.Histogram {
	var h metrics.Histogram
	if name != "" && archaius.Conf.Collect {
		h = expvar.NewHistogram(name, 1000, maxHistObservable, 1, []int{50, 99}...)
		sampleLock.Lock()
		if sampleMap == nil {
			sampleMap = make(map[metrics.Histogram][]int64)
		}
		sampleMap[h] = make([]int64, 0, sampleCount)
		sampleLock.Unlock()
		return h
	}
	return nil
}

func Measure(h metrics.Histogram, d time.Duration) {
	if h != nil && archaius.Conf.Collect {
		if d > maxHistObservable {
			h.Observe(int64(maxHistObservable))
		} else {
			h.Observe(int64(d))
		}
		sampleLock.Lock()
		s := sampleMap[h]
		if s != nil && len(s) < sampleCount {
			sampleMap[h] = append(s, int64(d))
		sampleLock.Unlock()
		}
	}
}
```

Median and 99%ile

Slice for first 500 values as samples for export to Guesstimate

Golang Guesstimate Interface

https://github.com/adrianco/goguesstimate


```
{
      "space": {
      "name": "gotest",
      "description": "Testing",
      "is_private": "true",
      "graph":  {
          "metrics": [
              {"id": "AB", "readableId": "AB", "name": "memcached", "location": {"row": 2, "column":4}},
              {"id": "AC", "readableId": "AC", "name": "memcached percent", "location": {"row": 2, "column":
3}},
              {"id": "AD", "readableId": "AD", "name": "staash cpu", "location": {"row": 3, "column":3}},
              {"id": "AE", "readableId": "AE", "name": "staash", "location": {"row": 3, "column":2}}
              ],
          "guesstimates": [
           {"metric": "AB", "input": null, "guesstimateType": "DATA", "data":
[119958,6066,13914,9595,6773,5867,2347,1333,9900,9404,13518,9021,7915,3733,10244,5461,12243,7931,9044,11706,
5706,22861,9022,48661,15158,28995,16885,9564,17915,6610,7080,7065,12992,35431,11910,11465,14455,25790,8339,9
991]},
              {"metric": "AC", "input": "40", "guesstimateType": "POINT"},
              {"metric": "AD", "input": "[1000,4000]", "guesstimateType": "LOGNORMAL"},
              {"metric": "AE", "input": "=100+((randomInt(0,100)>AC)?AB:AD)", "guesstimateType": "FUNCTION"}
          ]
      }
  }
}
```

See http://www.getguesstimate.com

```
% cd json_metrics; sh guesstimate.sh storage
```

Simplicity through symmetry

Symmetry
Invariants
Stable assertions
No special cases
Single purpose components

What's Next?

Serverless

Serverless Architectures

AWS Lambda starting to take off

Google Cloud Functions, Azure Functions on the way…

Open Source: IBM OpenWhisk, open-lambda.org, apex.run

Startup activity: iron.io , serverless.com, nano-lambda.com

With AWS Lambda compute resources are charged by the 100ms, not the hour

First 1M node.js executions/month are free

Monitorless Architecture

Monitorless Architecture

Monitorless Architecture

AWS Lambda Reference Arch

http://www.allthingsdistributed.com/2016/05/aws-lambda-serverless-reference-architectures.html

Serverless Programming Model

Event driven functions
Role based permissions
Whitelisted API based security
Good for simple single threaded code

Serverless Cost Efficiencies

100% useful work, no agents, overheads
100% utilization, no charge between requests
No need to size capacity for peak traffic
Anecdotal costs ~1% of conventional system
Ideal for low traffic, Corp IT, spiky workloads

Serverless Work in Progress

Tooling for ease of use
Multi-region HA/DR patterns
Debugging and testing frameworks
Monitoring tools vendor support
Tracing - see iopipe.com

DIY Serverless Operating Challenges

Startup latency
Execution overhead
Charging model
Capacity planning

Intelligent Monitoring

Sprinkle machine learning on everything!
Diagnosis and remediation - Instana, Stackstorm
Personalized role based user interfaces?
Canary signature analysis - Netflix Spinnaker & OpsMx.com
Event correlation & filtering - BigPanda

Monitoring Challenges

Too much new stuff
Too ephemeral
Too dumb

Teraservices

Terabyte Memory Directions

Engulf dataset in memory for analytics

Balanced config for memory intensive workloads

Replace high end systems at commodity cost point

Explore non-volatile memory implications

Terabyte Memory Options

Flash based Diablo 64/128/256GB DDR4 DIMM

Shipping now as volatile memory, future non-volatile

Intel 3D XPoint - new non-volatile technology on the way

Announced May 2016

AWS X1 Instance Type - over 2TB RAM for $14/hr

Easy availability should drive innovation

Diablo Memory1: Flash DIMM

MEMORY 1™

- ✓ **UP TO 256GB DDR4 MEMORY PER MODULE**
- ✓ **UP TO 4TB MEMORY IN 2 SOCKET SYSTEM**

- **NO CHANGES** to CPU or Server
- **NO CHANGES** to Operating System
- **NO CHANGES** to Applications

“We see the world as increasingly more complex and chaotic because we use inadequate concepts to explain it. When we understand something, we no longer see it as chaotic or complex.”

Jamshid Gharajedaghi - 2011
Systems Thinking: Managing Chaos and Complexity: A Platform for Designing Business Architecture

Learn More…

THE ART OF
BUSINESS VALUE
MARK SCHWARTZ

O'REILLY
Building
Microservices
DESIGNING FINE-GRAINED SYSTEMS
Sam Newman

AWS Lambda
A Guide to Serverless Microservices
By: Matthew Fuller

Q&A

Adrian Cockcroft @adrianco
http://slideshare.com/adriancockcroft
Technology Fellow - Battery Ventures

BV
Battery Ventures

See www.battery.com for a list of portfolio investments

BV Battery Ventures

Enterprise IT

Compute

NUTANIX
STRATO SCALE
mx mendix the app platform

Networking

BlueJeans
fastly
cumulus

Storage

DELPHIX
diablo technologies
Primary Data
Elastifile
Zerto

Data

platfora
CASK
bigpanda
SiSense
inter|ana

Operations & Management

AppDynamics
VividCortex
catchpoint
JFrog
Chef
OpsGenie
Dataloop.IO

Security

LIGHTCYBER CLOSING THE BREACH DETECTION GAP
RISKIQ
GuardiCore Securing SDN
VERA
CYVERA CYBER DEFENSE SOLUTIONS
Palo Alto Networks

Visit http://www.battery.com/our-companies/ for a full list of all portfolio companies in which all Battery Funds have invested.